



MULTITHREADING AND PARALLEL PROGRAMMING

Chapter 30- Liang

Roadmap

- Overview of multithreading (§30.2).
- Implementing the **Runnable** interface (§30.3).
- The **Thread** class (§30.3).
- Methods in the **Thread** class (§30.4).
- Thread pool (§30.6).
- Synchronize threads to avoid race conditions (§30.7).
- Synchronize threads using locks (§30.8).
- Thread communications using conditions on locks (§30.9–30.10).
- Restricting the number of accesses to a shared resource using semaphores (§30.12).
- Using the resource-ordering technique to avoid deadlocks (§30.13).
- Thread life cycle (§30.14).

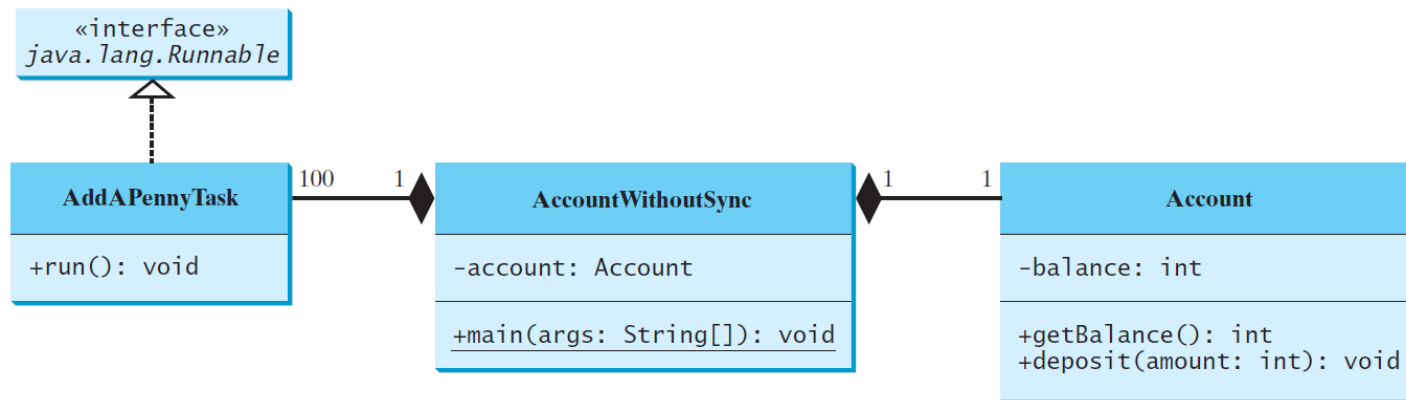
Thread Synchronization

- Thread **synchronization** is to **coordinate the execution of the dependent threads**.
- A shared resource may be corrupted if it is accessed simultaneously by multiple threads.
 - For example, two unsynchronized threads accessing the same bank account may cause conflict.

Step	balance	thread[i]	thread[j]
1	0	newBalance = bank.getBalance() + 1;	
2	0		newBalance = bank.getBalance() + 1;
3	1	bank.setBalance(newBalance);	
4	1		bank.setBalance(newBalance);

Example: Showing Resource Conflict

- Objective: Write a program that demonstrates the problem of resource conflict. Suppose that you create and launch one hundred threads, each of which adds a penny to an account. Assume that the account is initially empty.



```
Command Prompt
C:\book>java AccountWithoutSync
What is balance ? 5
C:\book>java AccountWithoutSync
What is balance ? 4
C:\book>java AccountWithoutSync
What is balance ? 7
C:\book>
```

The screenshot shows a Windows Command Prompt window titled "Command Prompt". It displays the execution of the `AccountWithoutSync` program three times. Each time, the program prints "What is balance ?" followed by a number. The first two runs show 5 and 4, while the third run shows 7, demonstrating the resource conflict where the balance is not updated correctly due to the lack of synchronization.

Race Condition

- What, then, caused the error in the example? Here is a possible scenario:

Step	Balance	Task 1	Task 2
1	0	<code>newBalance = balance + 1;</code>	
2	0		<code>newBalance = balance + 1;</code>
3	1	<code>balance = newBalance;</code>	
4	1		<code>balance = newBalance;</code>

- The effect of this scenario is that Task 1 did nothing, because in Step 4 Task 2 overrides Task 1's result.
- Obviously, the problem is that Task 1 and Task 2 are accessing a common resource in a way that causes conflict.
- This is a common problem known as a **race condition** in multithreaded programs.
- A class is said to be **thread-safe** if an **object of the class does not cause a race condition in the presence of multiple threads**.
- The Account class is not thread-safe.

The **synchronized** keyword

- To avoid race conditions, more than one thread must be prevented from simultaneously entering certain part of the program, known as critical region/section.
- The critical region is the entire deposit method.
- You can use the `synchronized` keyword to synchronize the method so that only one thread can access the method at a time.
- There are several ways to correct the problem,
 - one approach is to make Account thread-safe by adding the `synchronized` keyword in the deposit method in Line 45 as follows:

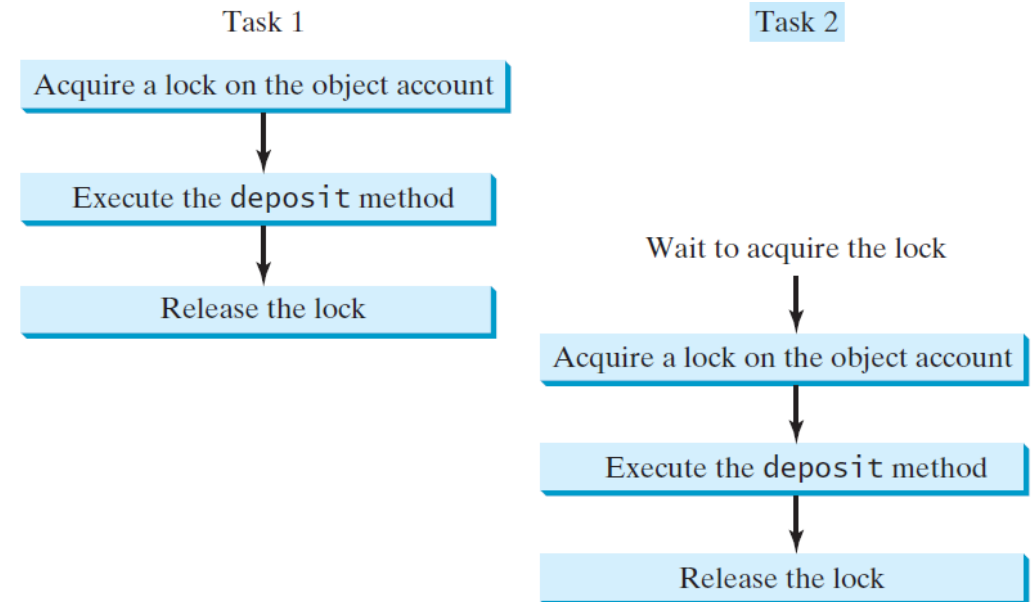
```
public synchronized void deposit(double amount)
```

Synchronizing Instance Methods and Static Methods

- A synchronized method acquires a lock before it executes.
 - A lock is a mechanism for exclusive use of a resource.
- In the case of an **instance method**, the **lock** is on the **object** for which the method was invoked.
 - If one thread invokes a synchronized instance method on an object, the lock of that object is acquired first, then the method is executed, and finally the lock is released.
 - Another thread invoking the same method of that object is blocked until the lock is released.
- In the case of a **static method**, the **lock** is on the **class**.
 - If one thread invokes a static method on an object, the lock of that class is acquired first, then the method is executed, and finally the lock is released.
 - Another thread invoking the same method of that class is blocked until the lock is released.

Synchronizing Tasks

- If the `deposit` method is synchronized, the earlier scenario cannot happen.
- If Task 2 starts to enter the method, and Task 1 is already in the method, Task 2 is blocked until Task 1 finishes the method.



Synchronizing Statements

- A synchronized statement can be used to acquire a lock on any object, not just *this* object, when executing a block of the code in a method.
- This block is referred to as a **synchronized block**. The general form of a synchronized statement is as follows:

```
synchronized (expr) {  
    statements;  
}
```

- The expression `expr` must evaluate to an object reference.
- If the **object is already locked by another thread, the thread is blocked until the lock is released.**
- When a lock is obtained on the object, the statements in the synchronized block are executed, and then the lock is released.

Synchronizing Statements vs. Methods

- Any synchronized instance method can be converted into a synchronized statement.
- Suppose that the following is a synchronized instance method:

```
public synchronized void xMethod() {  
    // method body  
}
```

- This method is equivalent to

```
public void xMethod() {  
    synchronized (this) {  
        // method body  
    }  
}
```

Example

- Four threads are accessing and modifying a shared Counter object named cntObj. What is the impact of invoking:
 - T4 invoking doubleValue() ?
 - T1, T2 invoking increment() and T3 invoking decrement() ?
 - T3 invoking increment() and T2, T4 invoking getValue() ?
- Address the possibility of race conditions.

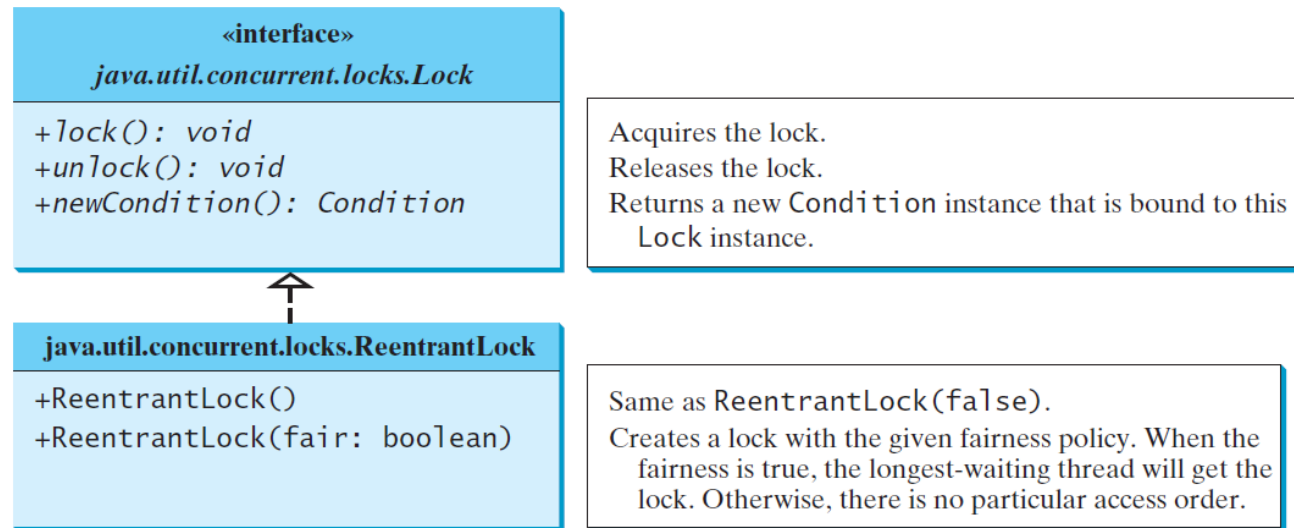
```
public class Counter {  
    private int count = 0;  
    public static synchronized void doubleValue() {  
        count = count * 2;  
    }  
    public synchronized void increment() {  
        count++;  
    }  
    public synchronized void decrement() {  
        count--;  
    }  
    public void update() {  
        count+=3;  
    }  
    public int getValue() {  
        return count;  
    }  
}
```

Synchronization Using Locks

- A synchronized instance method implicitly acquires a lock on the instance before it executes the method.
- Locks and conditions can be explicitly used to synchronize threads.
 - JDK 1.5 enables you to use locks explicitly.
- The new locking features are flexible and give you more control for coordinating threads.

Synchronization Using Locks

- A lock is an instance of the `Lock` interface, which declares the methods for acquiring and releasing locks.
- A lock may also use the `newCondition()` method to create any number of `Condition` objects, which can be used for thread communications.



Fairness Policy

- `ReentrantLock` is a concrete implementation of `Lock` for creating mutual exclusive locks.
- You can create a lock with the specified **fairness policy**.
 - True fairness policies guarantee the longest-wait thread to obtain the lock first.
 - False fairness policies grant a lock to a waiting thread without any access order.
- There are trade-offs between different types of locks used in multithreaded programming.
 - **Overall performance**: Measured by factors like throughput (requests processed per second) or latency (time taken to acquire the lock).
 - **Variance**: The spread of values around the average. Here, it refers to the variability in time it takes threads to acquire a lock.
 - **Starvation**: A situation where a thread waits indefinitely to acquire a lock due to other threads constantly taking it.

Fairness Policy

- Programs using fair locks may **have poorer overall performance** compared to those using the default (usually non-fair) setting.
 - This is because fair locks prioritize serving waiting threads in order, this can lead to situations where a busy thread constantly acquires the lock, delaying other waiting threads and potentially impacting overall throughput.
- Fair locks **offer smaller variances** in lock acquisition times.
 - This means wait times for threads are more predictable, unlike non-fair locks where a single thread might dominate access, causing some threads to wait significantly longer.
- Fair locks **prevent starvation**, which is a major concern with non-fair locks.
 - In non-fair scenarios, a thread might never get a chance to acquire the lock if other threads keep taking it, essentially starving it of access. Fair locks guarantee everyone eventually gets a turn, eliminating this risk.

Example: Using Locks

```
1 import java.util.concurrent.*;
2 import java.util.concurrent.locks.*;
3
4 public class AccountWithSyncUsingLock {
5     private static Account account = new Account();
6
7     public static void main(String[] args) {
8         ExecutorService executor = Executors.newCachedThreadPool();
9
10        // Create and launch 100 threads
11        for (int i = 0; i < 100; i++) {
12            executor.execute(new AddAPennyTask());
13        }
14
15        executor.shutdown();
16
17        // Wait until all tasks are finished
18        while (!executor.isTerminated()) {
19        }
20
21        System.out.println("What is balance ? " + account.getBalance());
22    }
23
24    // A thread for adding a penny to the account
25    public static class AddAPennyTask implements Runnable {
26        public void run() {
27            account.deposit(1);
28        }
29    }
30
```

```
31 // An inner class for account
32 public static class Account {
33     private static Lock lock = new ReentrantLock(); // Create a Lock
34     private int balance = 0;
35
36     public int getBalance() {
37         return balance;
38     }
39
40     public void deposit(int amount) {
41         lock.lock(); // Acquire the Lock
42
43         try {
44             int newBalance = balance + amount;
45
46             // This delay is deliberately added to magnify the
47             // data-corruption problem and make it easy to see.
48             Thread.sleep(5);
49
50             balance = newBalance;
51         }
52         catch (InterruptedException ex) {
53         }
54         finally {
55             lock.unlock(); // Release the Lock
56         }
57     }
58 }
59 }
```


Cooperation Among Threads

- The conditions can be used to facilitate communications among threads.
- A thread can specify what to do under a certain condition.
- Conditions are objects created by invoking the `newCondition()` method on a Lock object.
- Once a condition is created, you can use the methods `await()`, `signal()`, and `signalAll()` methods for thread communications

«interface» <i>java.util.concurrent.Condition</i>	
<i>+await(): void</i> <i>+signal(): void</i> <i>+signalAll(): Condition</i>	<i>Causes the current thread to wait until the condition is signaled.</i> <i>Wakes up one waiting thread.</i> <i>Wakes up all waiting threads.</i>

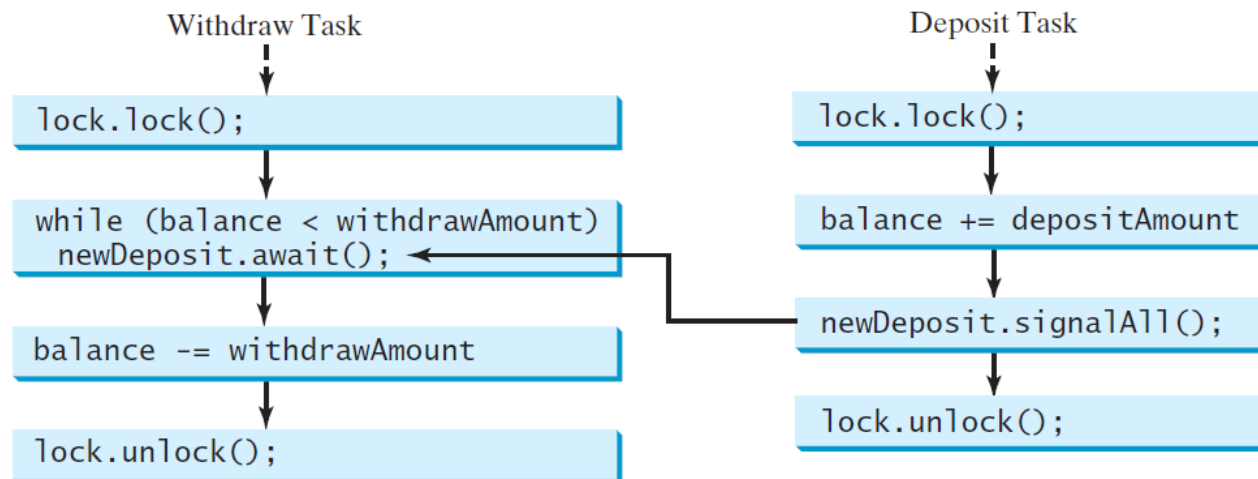
Cooperation Among Threads

- **How condition locks work?**

1. A thread **acquires the lock** associated with a shared resource.
2. The thread **checks** for the desired **condition** to be **true** (e.g., data available, task completed).
3. If the **condition is not true**, the thread uses **`await()`** to release the lock and voluntarily wait on the condition variable associated with the lock.
4. Another **thread fulfills** the condition by modifying the shared resource and calling **`signal()` or `signalAll()`** on the condition variable.
5. One or **all waiting threads are woken up** based on the **`signal()/signalAll()`** implementation and compete to reacquire the lock.
6. The **awakened thread re-checks** the condition and proceeds if it's true or waits again if not.

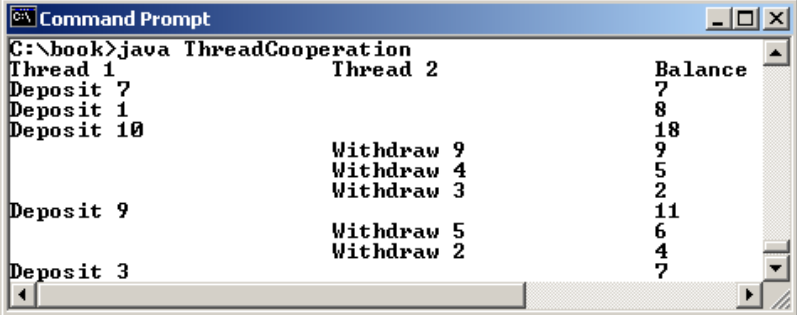
Cooperation Among Threads

- To **synchronize** the operations, use a lock with a condition: `newDeposit` (i.e., new deposit added to the account).
- If the `balance` is less than the amount to be withdrawn, the withdraw task will wait for the `newDeposit` condition.
- When the `deposit` task adds money to the account, the task signals the waiting withdraw task to try again.



Example: Thread Cooperation

- Write a program that demonstrates thread cooperation.
 - Suppose that you create and launch two threads, one deposits to an account, and the other withdraws from the same account.
 - The second thread must wait if the amount to be withdrawn is more than the current balance in the account.
 - Whenever new fund is deposited to the account, the first thread notifies the second thread to resume.
 - If the amount is still not enough for a withdrawal, the second thread must continue to wait for more fund in the account.
 - Assume the initial balance is 0 and the amount to deposit and to withdraw is randomly generated.



```
C:\book>java ThreadCooperation
Thread 1      Thread 2      Balance
Deposit 7      Withdraw 9      7
Deposit 1      Withdraw 4      8
Deposit 10     Withdraw 3      18
                Withdraw 5      9
                Withdraw 4      5
                Withdraw 3      2
Deposit 9      Withdraw 5      11
                Withdraw 2      6
Deposit 3      Withdraw 2      4
                Withdraw 2      7
```

Example: Thread Cooperation

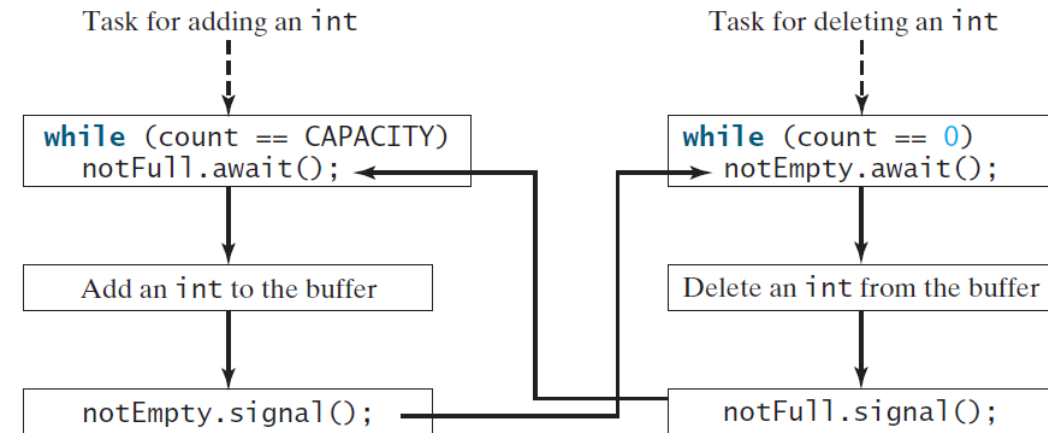
```
3 import java.util.concurrent.*;
4 import java.util.concurrent.locks.*;
5 public class ThreadCooperation {
6     private static Account account = new Account();
7     public static void main(String[] args) {
8         // Create a thread pool with two threads
9         ExecutorService executor = Executors.newFixedThreadPool(2);
10        executor.execute(new DepositTask());
11        executor.execute(new WithdrawTask());
12        executor.shutdown();
13        System.out.println("Thread 1\t\tThread 2\t\tBalance");
14    }
15    public static class DepositTask implements Runnable {
16        @Override // Keep adding an amount to the account
17        public void run() {
18            try { // Purposely delay it to let the withdraw method proceed
19                while (true) {
20                    account.deposit((int) (Math.random() * 10) + 1);
21                    Thread.sleep(1000);
22                }
23            } catch (InterruptedException ex) {
24                ex.printStackTrace();
25            }
26        }
27    }
28    public static class WithdrawTask implements Runnable {
29        @Override // Keep subtracting an amount from the account
30        public void run() {
31            while (true) {
32                account.withdraw((int) (Math.random() * 10) + 1);
33            }
34        }
35    }
```

```
36    private static class Account { // An inner class for account
37        private static Lock lock = new ReentrantLock(); // Create a new lock
38        // Create a condition
39        private static Condition newDeposit = lock.newCondition();
40        private int balance = 0;
41        public int getBalance() {
42            return balance;
43        }
44        public void withdraw(int amount) {
45            lock.lock(); // Acquire the lock
46            try {
47                while (balance < amount) {
48                    System.out.println("\t\t\tWait for a deposit");
49                    newDeposit.await();
50                }
51                balance -= amount;
52                System.out.println("\t\t\tWithdraw " + amount + "\t\t" + getBalance());
53            } catch (InterruptedException ex) {
54                ex.printStackTrace();
55            } finally {
56                lock.unlock(); // Release the lock
57            }
58        }
59        public void deposit(int amount) {
60            lock.lock(); // Acquire the lock
61            try {
62                balance += amount;
63                System.out.println("Deposit " + amount + "\t\t\t\t" + getBalance());
64                // Signal thread waiting on the condition
65                newDeposit.signalAll();
66            } finally {
67                lock.unlock(); // Release the lock
68            }
69        }
70    }
71 }
```

Case Study: Producer/Consumer

self-study

- Consider the classic Consumer/Producer example.
- Suppose you use a buffer to store integers.
 - The buffer size is limited.
- The buffer provides the method `write(int)` to add an int value to the buffer and the method `read()` to read and delete an int value from the buffer.
- To synchronize the operations, use a lock with two conditions: `notEmpty` (i.e., buffer is not empty) and `notFull` (i.e., buffer is not full).
- When a task adds an int to the buffer, if the buffer is full, the task will wait for the `notFull` condition.
- When a task deletes an int from the buffer, if the buffer is empty, the task will wait for the `notEmpty` condition.



Case Study: Producer/Consumer

self-study

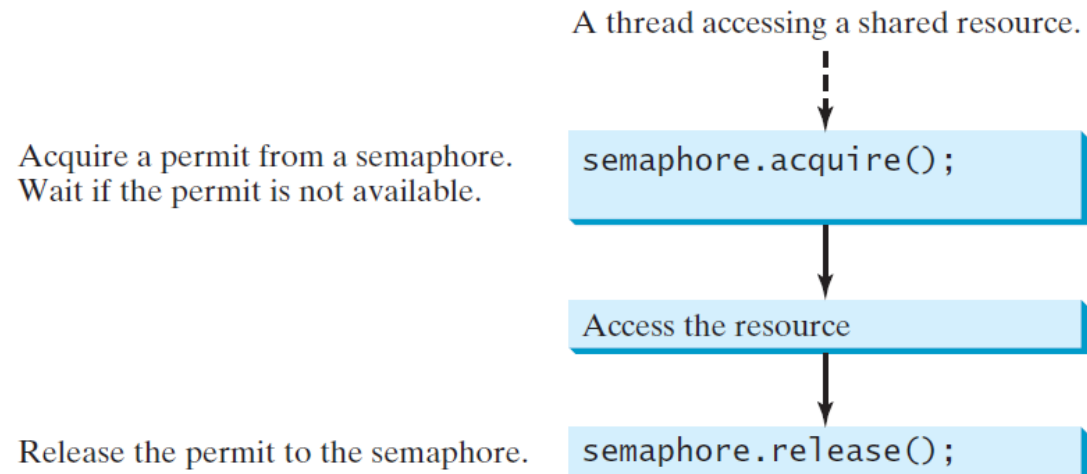
- Listing 30.8 presents the complete program. The program contains the Buffer class (lines 43-89) and two tasks for repeatedly producing and consuming numbers to and from the buffer (lines 15-41). The write(int) method (line 58) adds an integer to the buffer. The read() method (line 75) deletes and returns an integer from the buffer.
- For simplicity, the buffer is implemented using a linked list (lines 48-49). Two conditions notEmpty and notFull on the lock are created in lines 55-56. The conditions are bound to a lock. A lock must be acquired before a condition can be applied. If you use the wait() and notify() methods to rewrite this example, you have to designate two objects as monitors.

Remarks

- Once a thread invokes `await()` on a condition, the thread waits for a signal to resume.
- If you forget to call `signal()` or `signalAll()` on the condition, the thread will wait forever.
- A condition is created from a Lock object.
- To invoke the method (e.g., `await()`, `signal()`, and `signalAll()`), you must first own the lock.
- If you invoke these methods without acquiring the lock, an exception will be thrown.

Semaphores

- **Semaphores** can be used to **restrict the number of threads** that access a shared resource.
 - Before accessing the resource, a thread must **acquire a permit** from the semaphore.
 - After finishing with the resource, the thread must **return the permit** back to the semaphore.



Creating Semaphores

- To create a semaphore, you must specify the number of permits with an optional fairness policy.
- A task acquires a permit by invoking the semaphore's `acquire()` method and releases the permit by invoking the semaphore's `release()` method.
- Once a permit is acquired, the total number of available permits in a semaphore is reduced by 1.
- Once a permit is released, the total number of available permits in a semaphore is increased by 1.

`java.util.concurrent.Semaphore`

```
+Semaphore(numberOfPermits: int)
+Semaphore(numberOfPermits: int, fair:
  boolean)
+acquire(): void
+release(): void
```

Creates a semaphore with the specified number of permits. The fairness policy is false.

Creates a semaphore with the specified number of permits and the fairness policy.

Acquires a permit from this semaphore. If no permit is available, the thread is blocked until one is available.

Releases a permit back to the semaphore.

Creating Semaphores

- The code shows a semaphore with just one permit can be used to simulate a mutually exclusive lock.

```
1 // An inner class for Account
2 private static class Account {
3     // Create a semaphore
4     private static Semaphore semaphore = new Semaphore(1);
5     private int balance = 0;
6
7     public int getBalance() {
8         return balance;
9     }
10
11    public void deposit(int amount) {
12        try {
13            semaphore.acquire(); // Acquire a permit
14            int newBalance = balance + amount;
15
16            // This delay is deliberately added to magnify the
17            // data-corruption problem and make it easy to see
18            Thread.sleep(5);
19
20            balance = newBalance;
21        }
22        catch (InterruptedException ex) {
23        }
24        finally {
25            semaphore.release(); // Release a permit
26        }
27    }
28 }
```

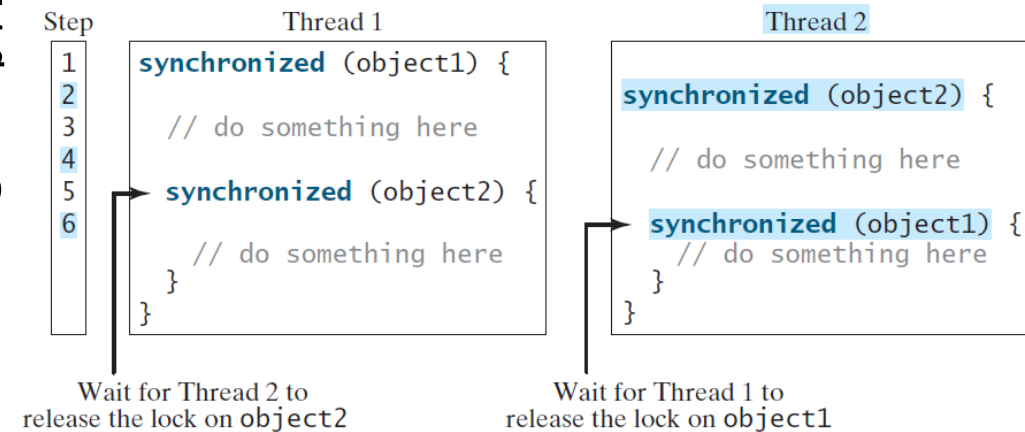
create a semaphore

acquire a permit

release a permit

Deadlock

- Sometimes two or more threads need to acquire the locks on several shared objects.
- This could cause **deadlock**, in which each thread has the lock on one of the objects and is waiting for the lock on the other object.
- Consider the scenario with two threads and two objects.
 - Thread 1 acquired a lock on object1 and Thread 2 acquired a lock on object2.
 - Now Thread 1 is waiting for the lock on object2 and Thread 2 for the lock on object1.
 - The two threads wait for each other to release the in order to get the lock, and neither can continue to run.

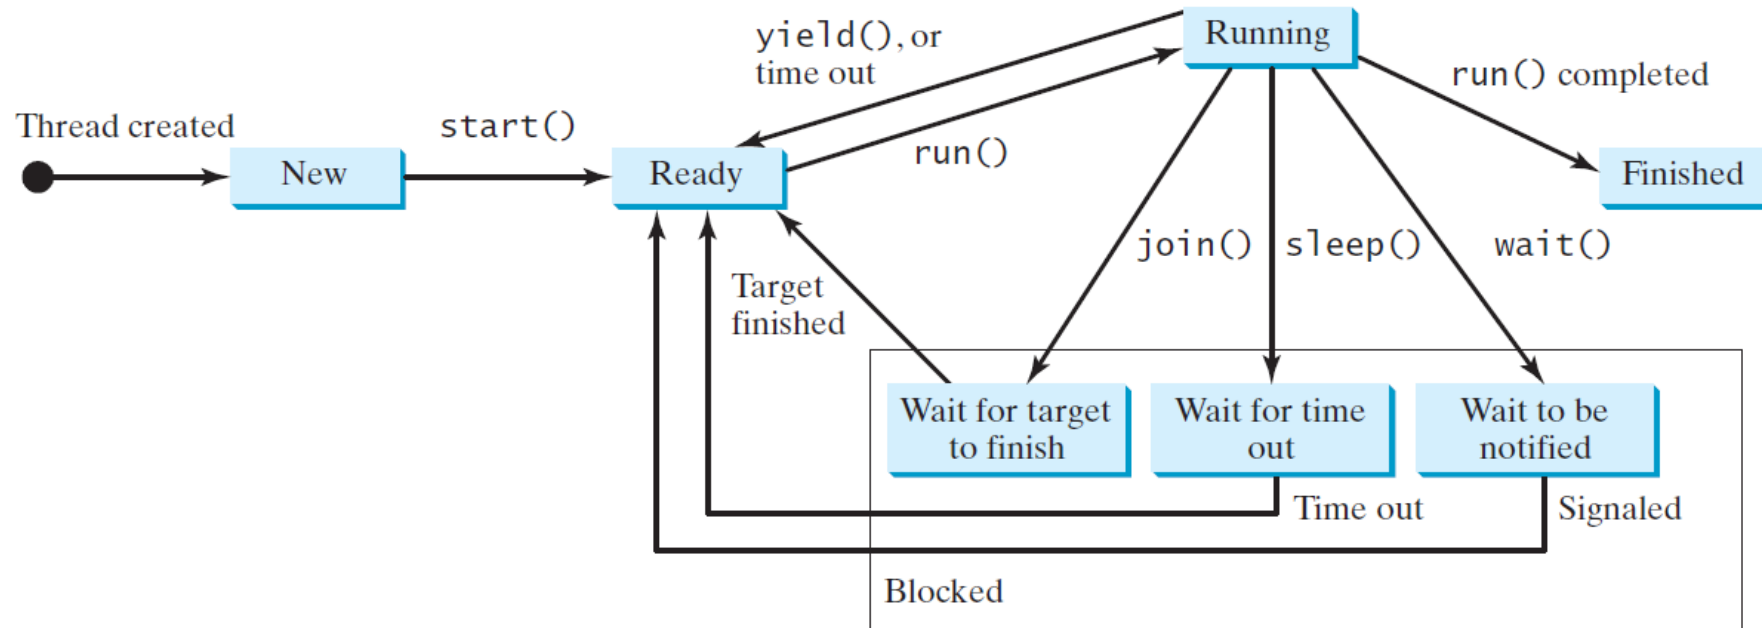


Preventing Deadlock

- Deadlock can be easily avoided by using a simple technique known as **resource ordering**.
- With this technique, you **assign an order on all the objects whose locks must be acquired** and ensure that each thread acquires the locks in that order.
- For the example, suppose the objects are ordered as object1 and object2.
 - Using the resource ordering technique, Thread 2 must acquire a lock on object1 first, then on object2.
 - Once Thread 1 acquired a lock on object1, Thread 2 must wait for a lock on object1.
 - Thread 1 will be able to acquire a lock on object2 and no deadlock would occur.

Thread States

- A thread can be in one of five states: New, Ready, Running, Blocked, or Finished.




Thread States

- A thread is in the **NEW** state after its creation using the Thread constructor but before calling its `start()` method. In this state, the thread hasn't allocated any resources and isn't eligible to run.
- After calling `start()`, the thread enters the **RUNNABLE** state. It's now eligible to run on a CPU core, waiting for its turn in the scheduling queue. Multiple threads can be in the **RUNNABLE** state simultaneously.
 - A ready thread is runnable but may not be running yet. The operating system must allocate CPU time to it.
- When the thread scheduler assigns a CPU core to a **RUNNABLE** thread, it transitions to the **RUNNING** state. It actively executes its code and consumes CPU resources. Only one thread can be in the **RUNNING** state on a single CPU core at a time.
 - A running thread can enter the Ready state if its CPU time expires, or its `yield()` method is called.
- A thread can enter the Blocked state (i.e., become inactive) for several reasons.
 - It may have invoked the `join(), sleep(), or wait()` method.
 - It may be waiting for an I/O operation to finish.

Thread States

- A thread enters the **BLOCKED** state when it encounters an event that prevents it from further execution, such as:
 - Waiting for I/O operations to complete (e.g., reading from a file).
 - Waiting for timeout.
 - Waiting for a notification from another thread using `wait()` or `join()`.
- A **BLOCKED** thread may be reactivated when the action caused the inactivation is reversed.
- The `isAlive()` method is used to find out the state of a thread.
 - It returns true if a thread is in the Ready, Blocked, or Running state.
 - It returns false if a thread is New and has not started or if it is Finished.
- The `interrupt()` method interrupts a thread in the following way:
 - If a thread is currently in the Ready or Running state, its interrupted flag is set.
 - If a thread is currently Blocked, it is awakened and enters the Ready state, and a `java.lang.InterruptedException` is thrown.
- A thread reaches the **FINISHED** state when it finishes executing its `run` method or explicitly throws an uncaught exception. It has released all its resources and can't be restarted.



ANY QUESTIONS?